

# speech-detection-assignment

August 2, 2023

## 1 Spoken Digit Recognition

In this notebook, You will do Spoken Digit Recognition.

Input - speech signal, output - digit number

It contains

1. Reading the dataset. and Preprocess the data set. Detailed instructions are given below.
2. You have to write the code in the same cell which contains the instruction.
3. Training the LSTM with RAW data
4. Converting to spectrogram and Training the LSTM network
5. Creating the augmented data and doing step 2 and 3 again.

```
[ ]: from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

```
[ ]: !unzip "/content/drive/MyDrive/projects_DS_ML/Spoken_Digit_recognition/
↳ recordings.zip"
```

```
[ ]: import numpy as np
import pandas as pd
import librosa
import os
from sklearn.model_selection import train_test_split
import matplotlib.pyplot as plt
import seaborn as sns
from tensorflow.keras.preprocessing.sequence import pad_sequences
```

We shared recordings.zip, please unzip those.

```
[ ]: ## defining a path for the dataset
path = "/content/recordings"
all_files = os.listdir(path)
```

```
[ ]: int(all_files[0][0])
```

```
[ ]: 0
```

Grader function 1

Create a dataframe(name=df\_audio) with two columns(path, label).

You can get the label from the first letter of name.

Eg: 0\_jackson\_0 -> 0

0\_jackson\_43 -> 0

## 1.1 Creating dataframe

```
[ ]: label= [int(i[0]) for i in all_files]
df_audio= pd.DataFrame(list(zip(all_files, label)),
                        columns=['path', 'label'])
```

```
[ ]: df_audio
```

```
[ ]:
      path  label
0    0_theo_29.wav    0
1    5_theo_25.wav    5
2    7_theo_6.wav    7
3    8_theo_31.wav    8
4    7_jackson_23.wav    7
...
1995  6_nicolas_37.wav    6
1996  7_theo_24.wav    7
1997  8_yweweler_48.wav    8
1998  0_jackson_4.wav    0
1999  6_theo_9.wav    6
```

[2000 rows x 2 columns]

```
[ ]: #info
df_audio.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2000 entries, 0 to 1999
Data columns (total 2 columns):
#   Column  Non-Null Count  Dtype
---  -
0   path    2000 non-null       object
1   label   2000 non-null       int64
dtypes: int64(1), object(1)
memory usage: 31.4+ KB
```

```
[ ]: from sklearn.utils import shuffle
df_audio = shuffle(df_audio, random_state=33)#don't change the random state
```

```
[ ]: df_audio['path'].shape
```

```
[ ]: (2000,)
```

```
[ ]: y=df_audio["label"]    ##this gives x and y as pandas series
X=df_audio["path"]
X_train, X_test, y_train, y_test=train_test_split(X, y, test_size=0.30,
↳stratify=y, random_state=33)
```

```
[ ]: type(X_train)    ##pandas series have value_counts functions
```

```
[ ]: pandas.core.series.Series
```

we will create two functions using which a wav file can be converted to spectrogram

```
[ ]: sample_rate = 22050
def load_wav(x, get_duration=True):
    '''This will return the array values of audio with sampling rate of 22050,
↳and Duration'''
    #loading the wav file with sampling rate of 22050
    final = path + '/' + x
    samples, sample_rate = librosa.load(final, sr=22050)
    if get_duration:
        duration = librosa.get_duration(filename = final)
        return [samples, duration]
    else:
        return samples
```

```
[ ]: def preprocess_wav(arr):
    samples = []
    durations = []
    for i in arr.values:
        sample, duration = load_wav(i)
        samples.append(sample)
        durations.append(duration)
    dataframe_processed = pd.DataFrame(list(zip(samples, durations)),columns=
↳['raw_data', 'duration'])
    return dataframe_processed
```

```
[ ]: # train_samples , train_durations = preprocess_wav(X_train)
preprocessed_train = preprocess_wav(X_train)
```

<ipython-input-13-bfb723a3eac9>:8: FutureWarning: get\_duration() keyword argument 'filename' has been renamed to 'path' in version 0.10.0.

This alias will be removed in version 1.0.

```
duration = librosa.get_duration(filename = final)
```

```
[ ]: preprocessed_test = preprocess_wav(X_test)
```

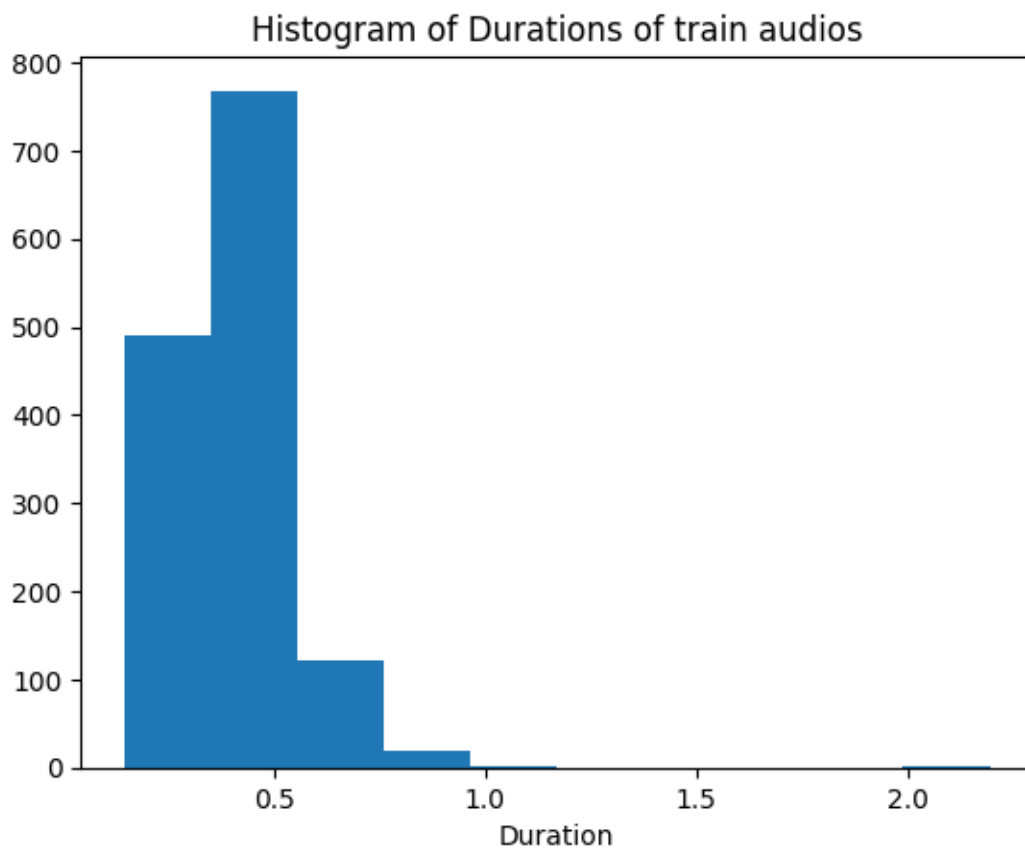
<ipython-input-13-bfb723a3eac9>:8: FutureWarning: get\_duration() keyword

argument 'filename' has been renamed to 'path' in version 0.10.0.  
This alias will be removed in version 1.0.  
duration = librosa.get\_duration(filename = final)


```
[ ]: print(np.percentile(preprocessed_train['duration'].values, 50))
```

0.390875

```
[ ]: #plot the histogram of the duration for trian
plt.hist(preprocessed_train.duration)
plt.xlabel('Duration')
plt.title('Histogram of Durations of train audios')
plt.show()
```



```
[ ]: #plot the histogram of the duration for test
plt.hist(preprocessed_test.duration)
plt.xlabel('Duration')
plt.title('Histogram of Durations of test audios')
plt.show()
```



```
[ ]: #print 0 to 100 percentile values with step size of 10 for train data duration.  
for i in range(0,110,10):  
    print(i,"th percentile is ",np.percentile(preprocessed_train['duration'].  
↪values, i))
```

```
0 th percentile is  0.1435  
10 th percentile is  0.2627  
20 th percentile is  0.30275  
30 th percentile is  0.3360875  
40 th percentile is  0.3615  
50 th percentile is  0.390875  
60 th percentile is  0.417575  
70 th percentile is  0.4462  
80 th percentile is  0.482625  
90 th percentile is  0.5554  
100 th percentile is  2.195875
```

```
[ ]: ##print 90 to 100 percentile values with step size of 1.  
for i in range(90,101):
```

```
print(i,"th percentile is ",np.percentile(preprocessed_train['duration'].
↪values, i))
```

```
90 th percentile is 0.5554
91 th percentile is 0.5707725
92 th percentile is 0.5811350000000001
93 th percentile is 0.5915062500000003
94 th percentile is 0.61042
95 th percentile is 0.6230937499999999
96 th percentile is 0.6435949999999999
97 th percentile is 0.6647649999999999
98 th percentile is 0.7050174999999999
99 th percentile is 0.8072624999999999
100 th percentile is 2.195875
```

Based on our analysis 99 percentile values are less than 0.8sec so we will limit maximum length of X\_train\_processed and X\_test\_processed to 0.8 sec. It is similar to pad\_sequence for a text dataset.

While loading the audio files, we are using sampling rate of 22050 so one sec will give array of length 22050. so, our maximum length is  $0.8 \times 22050 = 17640$  Pad with Zero if length of sequence is less than 17640 else Truncate the number.

Also create a masking vector for train and test.

masking vector value = 1 if it is real value, 0 if it is pad value. Masking vector data type must be bool.

```
[ ]: preprocessed_train['raw_data'][2].shape    ###every row has different number of
↪values in array depending on duration
```

```
[ ]: (9992,)
```

```
[ ]: max_length = 17640
```

```
[ ]: ## as discussed above, Pad with Zero if length of sequence is less than 17640
↪else Truncate the number.
## save in the X_train_pad_seq, X_test_pad_seq
## also Create masking vector X_train_mask, X_test_mask

## all the X_train_pad_seq, X_test_pad_seq, X_train_mask, X_test_mask will be
↪numpy arrays mask vector dtype must be bool.
def pad_mask(df, max_length):
    new_padded=[]
    new_masks = []
    for i in range(df['raw_data'].shape[0]):
        length = df['raw_data'][i].shape[0]
        if length <= max_length:
            new_add = [0] * (max_length - length)
```

```

new_one = np.concatenate((df['raw_data'][i],new_add), axis=0)
mask = [True]* length + [False]* (max_length - length)

else:
    new_one = df['raw_data'][i][: max_length]    ##https://www.geeksforgeeks.
    ↪org/python-remove-last-k-elements-of-list/#:~:
    ↪text=%23%20initializing%20list%C2%A0,str(res))
    mask = [True]*max_length
    # new_one = np.concatenate((X_train_processed['raw_data'][i],new_add),
    ↪axis=0)
    new_padded.append(new_one)
    new_masks.append(mask)
return np.array(new_padded), np.array(new_masks)

```

```

[ ]: X_train_pad_seq, X_train_mask = pad_mask(preprocessed_train, max_length)
X_test_pad_seq, X_test_mask = pad_mask(preprocessed_test, max_length)

```

```

[ ]: X_train_mask[10]

```

```

[ ]: array([ True,  True,  True, ..., False, False, False])

```

saving the numpy arrays to disc to use them directly for further assessment

```

[ ]: # np.save('X_train_pad_seq.npy', X_train_pad_seq )
# np.save('X_test_pad_seq.npy', X_test_pad_seq)
# np.save('X_train_mask.npy', X_train_mask )
# np.save('X_test_mask.npy', X_test_mask)

```

Loading the arrays

```

[ ]: # import numpy as np

# X_train_pad_seq = np.load('/content/drive/MyDrive/projects_DS_ML/
    ↪Spoken_Digit_recognition/X_train_pad_seq.npy')
# X_test_pad_seq = np.load('/content/drive/MyDrive/projects_DS_ML/
    ↪Spoken_Digit_recognition/X_test_pad_seq.npy')
# X_train_mask = np.load('/content/drive/MyDrive/projects_DS_ML/
    ↪Spoken_Digit_recognition/X_train_mask.npy')
# X_test_mask = np.load('/content/drive/MyDrive/projects_DS_ML/
    ↪Spoken_Digit_recognition/X_test_mask.npy')

```

### 1.1.1 1. Giving Raw data directly.

Now we have

Train data: X\_train\_pad\_seq, X\_train\_mask and y\_train

Test data: X\_test\_pad\_seq, X\_test\_mask and y\_test

We will create a LSTM model which takes this input.

Task:

1. Create an LSTM network which takes “X\_train\_pad\_seq” as input, “X\_train\_mask” as mask input. You can use any number of LSTM cells. Please read LSTM documentation([https://www.tensorflow.org/api\\_docs/python/tf/keras/layers/LSTM](https://www.tensorflow.org/api_docs/python/tf/keras/layers/LSTM)) in tensorflow to know more about mask and also [https://www.tensorflow.org/guide/keras/masking\\_and\\_padding](https://www.tensorflow.org/guide/keras/masking_and_padding)
2. Get the final output of the LSTM and give it to Dense layer of any size and then give it to Dense layer of size 10(because we have 10 outputs) and then compile with the sparse categorical cross entropy( because we are not converting it to one hot vectors). Also check the datatype of class labels(y\_values) and make sure that you convert your class labels to integer datatype before fitting in the model.
3. While defining your model make sure that you pass both the input layer and mask input layer as input to lstm layer as follows
4. Use tensorboard to plot the graphs of loss and metric(use custom micro F1 score as metric) and histograms of gradients. You can write your code for computing F1 score using this link
5. make sure that it won't overfit.
6. You are free to include any regularization

```
[ ]: from tensorflow.keras.layers import Input, LSTM, Dense
from tensorflow.keras.models import Model
import tensorflow as tf
from tensorflow.keras import layers
```

```
[ ]: import tensorflow as tf
import numpy as np
from sklearn.metrics import f1_score

class Metrics(tf.keras.callbacks.Callback):
    def __init__(self, validation_data):
        super().__init__()
        self.x_test = validation_data[0]
        self.y_test = validation_data[1]

    def on_epoch_end(self, epoch, logs=None):
        val_predict = np.asarray(self.model.predict(self.x_test))
        val_label = np.argmax(val_predict, axis=1)
        val_targ = self.y_test
        val_f1 = f1_score(val_targ, val_label, average='micro')
        print("val_F1_score:", val_f1)
```

```
[ ]: padded = Input(shape = (17640,1))
masked = Input(shape = (17640), dtype = 'bool')
layer = LSTM(25)(padded, mask = masked)
layer = Dense(50, activation = 'relu', kernel_initializer = 'he_normal')(layer)
```



```

out = Dense(10, activation = 'softmax', kernel_initializer =
↳'glorot_normal')(layer)

model1 = Model(inputs = [padded, masked], outputs = out)
#printing the summary of model
model1.summary()

```

Model: "model"

```

-----
Layer (type)                 Output Shape          Param #   Connected to
=====
input_1 (InputLayer)         [(None, 17640, 1)]    0         []
input_2 (InputLayer)         [(None, 17640)]       0         []
lstm (LSTM)                  (None, 25)            2700      ['input_1[0][0]',
'input_2[0][0]']
dense (Dense)                (None, 50)            1300      ['lstm[0][0]']
dense_1 (Dense)              (None, 10)            510       ['dense[0][0]']
=====
Total params: 4,510
Trainable params: 4,510
Non-trainable params: 0
-----

```

```

[ ]: import datetime
logdir = os.path.join('logs', datetime.datetime.now().strftime('%Y%m%d-%H%M%S'))

```

```

[ ]: tf.keras.backend.clear_session()
model1.compile(optimizer = tf.keras.optimizers.Adam(0.001), loss =
↳'sparse_categorical_crossentropy')

custom_calls = [tf.keras.callbacks.TensorBoard(logdir, histogram_freq = 1,
↳write_graph = True),
Metrics([[X_test_pad_seq, X_test_mask], y_test)),
tf.keras.callbacks.EarlyStopping(monitor = 'val_loss', min_delta = 0.1,
↳patience = 3)]

```

```
model1.fit([X_train_pad_seq, X_train_mask], y_train, validation_data =
↳([X_test_pad_seq, X_test_mask], y_test),
        batch_size = 16, epochs = 10, callbacks = custom_calls)
```

```
Epoch 1/10
19/19 [=====] - 4s 164ms/step
val_F1_score: 0.10499999999999998
88/88 [=====] - 51s 448ms/step - loss: 2.3035 -
val_loss: 2.3026
Epoch 2/10
19/19 [=====] - 4s 195ms/step
val_F1_score: 0.10000000000000002
88/88 [=====] - 40s 458ms/step - loss: 2.3035 -
val_loss: 2.3026
Epoch 3/10
19/19 [=====] - 3s 133ms/step
val_F1_score: 0.10000000000000002
88/88 [=====] - 37s 426ms/step - loss: 2.3034 -
val_loss: 2.3026
Epoch 4/10
19/19 [=====] - 3s 134ms/step
val_F1_score: 0.10000000000000002
88/88 [=====] - 34s 392ms/step - loss: 2.3030 -
val_loss: 2.3026
```

```
[ ]: <keras.callbacks.History at 0x7a9a80b0ec20>
```

```
[ ]: %load_ext tensorboard
      %tensorboard --logdir logs/
```

```
<IPython.core.display.Javascript object>
```

### 1.1.2 2. Converting into spectrogram and giving spectrogram data as input

We can use librosa to convert raw data into spectrogram. A spectrogram shows the features in a two-dimensional representation with the intensity of a frequency at a point in time i.e we are converting Time domain to frequency domain. you can read more about this in <https://pnsn.org/spectrograms/what-is-a-spectrogram>

```
[ ]: X_train_pad_seq.shape
```

```
[ ]: (1400, 17640)
```

```
[ ]: def convert_to_spectrogram(raw_data):
      '''converting to spectrogram'''
      spectrum = librosa.feature.melspectrogram(y=raw_data, sr=sample_rate,
↳n_mels=64)
      logmel_spectrum = librosa.power_to_db(S=spectrum, ref=np.max)
```

```
return logmel_spectrum
```

```
[ ]: def create_spectrogram(arr):  
    spect_list = []  
    for i in arr:  
        spect = convert_to_spectrogram(i)  
        spect_list.append(spect)  
    return(np.array(spect_list))
```

```
[ ]: X_train_spectrogram = create_spectrogram(X_train_pad_seq)
```

```
[ ]: X_test_spectrogram = create_spectrogram(X_test_pad_seq)
```

```
[ ]: X_test_spectrogram.shape
```

```
[ ]: (600, 64, 35)
```

shape we get is 1400, 64, 35 here 64 is n\_mels and 35 is no of frames There is a formula to calculate no of frames

Now we have

Train data: X\_train\_spectrogram and y\_train

Test data: X\_test\_spectrogram and y\_test

We will create a LSTM model which takes this input.

Task:

1. Create an LSTM network which takes “X\_train\_spectrogram” as input and has to return output at every time step.
2. Average the output of every time step and give this to the Dense layer of any size. (ex: Output from LSTM will be (None, time\_steps, features) average the output of every time step i.e, you should get (None, time\_steps) and then pass to dense layer )
3. give the above output to Dense layer of size 10( output layer) and train the network with sparse categorical cross entropy.
4. Use tensorboard to plot the graphs of loss and metric(use custom micro F1 score as metric) and histograms of gradients. You can write your code for computing F1 score using this link
5. make sure that it won't overfit.
6. You are free to include any regularization

```
[ ]: # write the architecture of the model  
#print model.summary and make sure that it is following point 2 mentioned above  
input = Input(shape = (64,35,))  
layer = LSTM(256, return_sequences = True)(input)  
layer = tf.math.reduce_mean(layer, axis = -1)  
layer = Dense(256, activation = 'relu', kernel_initializer = 'he_normal')(layer)  
layer = Dense(128, activation = 'relu', kernel_initializer =  
    ⇨ 'glorot_normal')(layer)
```

```

output = Dense(10, activation = 'softmax', kernel_initializer =_
↳'glorot_normal')(layer)

model2 = Model(inputs = input, outputs = output)
#printing the model summary
model2.summary()

```

Model: "model"

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 64, 35)]	0
lstm (LSTM)	(None, 64, 256)	299008
tf.math.reduce_mean (TFOpLa mbda)	(None, 64)	0
dense (Dense)	(None, 256)	16640
dense_1 (Dense)	(None, 128)	32896
dense_2 (Dense)	(None, 10)	1290

=====  
 Total params: 349,834  
 Trainable params: 349,834  
 Non-trainable params: 0  
 =====

1. input layer = batch\_size \* 64\*35
2. in 2nd layer, 64 \*128 i.e. each row passes with 128 cells
3. as return seq is true we get output 64128 otherwise it would be 641
4. now we use global pooling i.e. mean value of 128 cells so output = 64\*1
- 5.

```
[ ]: !rm -rf ./logs/
```

```

[ ]: #compile and fit your model.
#model2.fit([X_train_spectrogram],y_train_int,.....)

model2.compile(optimizer = tf.keras.optimizers.Adam(0.001), loss =_
↳'sparse_categorical_crossentropy')
callback_list = [tf.keras.callbacks.TensorBoard(logdir, histogram_freq = 1,_
↳write_graph = True),
Metrics([X_test_spectrogram, y_test]),

```

```
tf.keras.callbacks.EarlyStopping(monitor = 'val_loss', min_delta = 0.001,
    patience = 5)]

model2.fit(X_train_spectrogram, y_train, validation_data = (X_test_spectrogram,
    y_test),
    batch_size = 16, epochs = 50, callbacks = callback_list)
```

Epoch 1/50

1/88 [...] - ETA: 4:17 - loss: 2.3034

WARNING:tensorflow:Callback method `on\_train\_batch\_end` is slow compared to the batch time (batch time: 0.0068s vs `on\_train\_batch\_end` time: 0.0112s). Check your callbacks.

19/19 [=====] - 0s 3ms/step

val\_F1\_score: 0.49833333333333335

88/88 [=====] - 5s 23ms/step - loss: 1.9386 - val\_loss: 1.3828

Epoch 2/50

19/19 [=====] - 0s 3ms/step

val\_F1\_score: 0.6116666666666667

88/88 [=====] - 1s 10ms/step - loss: 1.2669 - val\_loss: 1.0695

Epoch 3/50

19/19 [=====] - 0s 3ms/step

val\_F1\_score: 0.6833333333333333

88/88 [=====] - 1s 10ms/step - loss: 0.9862 - val\_loss: 0.8880

Epoch 4/50

19/19 [=====] - 0s 3ms/step

val\_F1\_score: 0.6816666666666666

88/88 [=====] - 1s 10ms/step - loss: 0.8373 - val\_loss: 0.7923

Epoch 5/50

19/19 [=====] - 0s 4ms/step

val\_F1\_score: 0.805

88/88 [=====] - 1s 16ms/step - loss: 0.7399 - val\_loss: 0.6203

Epoch 6/50

19/19 [=====] - 0s 5ms/step

val\_F1\_score: 0.8000000000000002

88/88 [=====] - 1s 15ms/step - loss: 0.6340 - val\_loss: 0.5824

Epoch 7/50

19/19 [=====] - 0s 3ms/step

val\_F1\_score: 0.8066666666666665

88/88 [=====] - 1s 16ms/step - loss: 0.5646 - val\_loss: 0.5453

```

Epoch 8/50
19/19 [=====] - 0s 3ms/step
val_F1_score: 0.8483333333333335
88/88 [=====] - 1s 11ms/step - loss: 0.5429 - val_loss:
0.4749
Epoch 9/50
19/19 [=====] - 0s 3ms/step
val_F1_score: 0.8266666666666667
88/88 [=====] - 1s 11ms/step - loss: 0.4822 - val_loss:
0.4744
Epoch 10/50
19/19 [=====] - 0s 3ms/step
val_F1_score: 0.8483333333333335
88/88 [=====] - 1s 11ms/step - loss: 0.4657 - val_loss:
0.4788
Epoch 11/50
19/19 [=====] - 0s 3ms/step
val_F1_score: 0.865
88/88 [=====] - 1s 11ms/step - loss: 0.4089 - val_loss:
0.3832
Epoch 12/50
19/19 [=====] - 0s 3ms/step
val_F1_score: 0.875
88/88 [=====] - 1s 11ms/step - loss: 0.3916 - val_loss:
0.3887
Epoch 13/50
19/19 [=====] - 0s 3ms/step
val_F1_score: 0.8866666666666667
88/88 [=====] - 1s 11ms/step - loss: 0.3758 - val_loss:
0.3759
Epoch 14/50
19/19 [=====] - 0s 3ms/step
val_F1_score: 0.8716666666666667
88/88 [=====] - 1s 10ms/step - loss: 0.3530 - val_loss:
0.3797
Epoch 15/50
19/19 [=====] - 0s 3ms/step
val_F1_score: 0.87
88/88 [=====] - 1s 10ms/step - loss: 0.3365 - val_loss:
0.3738
Epoch 16/50
19/19 [=====] - 0s 3ms/step
val_F1_score: 0.8883333333333333
88/88 [=====] - 1s 11ms/step - loss: 0.3188 - val_loss:
0.3603
Epoch 17/50
19/19 [=====] - 0s 3ms/step
val_F1_score: 0.8666666666666667

```

```

88/88 [=====] - 1s 10ms/step - loss: 0.3258 - val_loss:
0.3783
Epoch 18/50
19/19 [=====] - 0s 4ms/step
val_F1_score: 0.8933333333333333
88/88 [=====] - 1s 14ms/step - loss: 0.3111 - val_loss:
0.3271
Epoch 19/50
19/19 [=====] - 0s 4ms/step
val_F1_score: 0.9016666666666667
88/88 [=====] - 1s 16ms/step - loss: 0.2869 - val_loss:
0.3048
Epoch 20/50
19/19 [=====] - 0s 3ms/step
val_F1_score: 0.8866666666666667
88/88 [=====] - 1s 16ms/step - loss: 0.2873 - val_loss:
0.3137
Epoch 21/50
19/19 [=====] - 0s 3ms/step
val_F1_score: 0.9183333333333333
88/88 [=====] - 1s 11ms/step - loss: 0.2583 - val_loss:
0.2613
Epoch 22/50
19/19 [=====] - 0s 3ms/step
val_F1_score: 0.8783333333333333
88/88 [=====] - 1s 11ms/step - loss: 0.2644 - val_loss:
0.3548
Epoch 23/50
19/19 [=====] - 0s 3ms/step
val_F1_score: 0.9033333333333333
88/88 [=====] - 1s 10ms/step - loss: 0.2579 - val_loss:
0.3052
Epoch 24/50
19/19 [=====] - 0s 3ms/step
val_F1_score: 0.8933333333333333
88/88 [=====] - 1s 10ms/step - loss: 0.2492 - val_loss:
0.3203
Epoch 25/50
19/19 [=====] - 0s 3ms/step
val_F1_score: 0.8916666666666667
88/88 [=====] - 1s 11ms/step - loss: 0.2588 - val_loss:
0.3074
Epoch 26/50
19/19 [=====] - 0s 3ms/step
val_F1_score: 0.9116666666666666
88/88 [=====] - 1s 11ms/step - loss: 0.2740 - val_loss:
0.2633

```

```
[ ]: <keras.callbacks.History at 0x7a9a4c70cc70>
```

```
[ ]: %tensorboard --logdir logs/
```

```
Reusing TensorBoard on port 6006 (pid 3892), started 0:02:16 ago. (Use '!kill_
↵3892' to kill it.)
```

```
<IPython.core.display.Javascript object>
```

### 1.1.3 3. Data augmentation with raw features

Till now we have done with 2000 samples only. It is very less data. We are giving the process of generating augmented data below.

There are two types of augmentation: 1. time stretching - Time stretching either increases or decreases the length of the file. For time stretching we move the file 30% faster or slower 2. pitch shifting - pitch shifting moves the frequencies higher or lower. For pitch shifting we shift up or down one half-step.

```
[ ]: ## generating augmented data.
def generate_augmented_data(file_path):
    augmented_data = []
    samples = load_wav(file_path,get_duration=False)
    for time_value in [0.7, 1, 1.3]: ## 30% change
        for pitch_value in [-1, 0, 1]:
            time_stretch_data = librosa.effects.time_stretch(samples,↵
↵rate=time_value)
            final_data = librosa.effects.pitch_shift(time_stretch_data,↵
↵sr=sample_rate, n_steps=pitch_value)
            augmented_data.append(final_data)
    return augmented_data
```

```
[ ]: temp_path = df_audio.iloc[0].path
aug_temp = generate_augmented_data(temp_path)
```

```
[ ]: len(aug_temp)
```

```
[ ]: 9
```

9 datapoints for one point

## 1.2 Follow the steps

1. Split data 'df\_audio' into train and test (80-20 split)
2. We have 2000 data points(1600 train points, 400 test points)

```
[ ]: X_train, X_test, y_train,↵
↵y_test=train_test_split(df_audio['path'],df_audio['label'],random_state=45,test_size=0.
↵2,stratify=df_audio['label'])
```



```
[ ]: X_train.shape
```

```
[ ]: (1600,)
```

3. Do augmentation only on X\_train, pass each point of X\_train to generate\_augmented\_data function. After augmentation we will get 14400 train points. Make sure that you are augmenting the corresponding class labels (y\_train) also.
4. Preprocess your X\_test using load\_wav function.
5. Convert the augmented\_train\_data and test\_data to numpy arrays.
6. Perform padding and masking on augmented\_train\_data and test\_data.
7. After padding define the model similar to model 1 and fit the data

```
[ ]: def create_aug_samples (arr,label):  
    aug_data = []  
    aug_label = []  
    for pt , label in zip(arr.values, label.values):  
  
        new_data = generate_augmented_data(pt)  
        aug_data.extend(new_data)  
        aug_label.extend([label]*9)  
    # return aug_data , aug_label  
    X_train_processed = pd.DataFrame(list(zip(aug_data, aug_label)), columns=  
    ↪=['raw_data', 'label'])  
    return X_train_processed
```

```
[ ]: X_train_processed = create_aug_samples(X_train, y_train)
```

```
[ ]: test_data = []  
for i in X_test.values:  
    test_data.append(load_wav(i, get_duration = False))  
X_test_processed = pd.DataFrame({'raw_data' : test_data, 'label' : y_test.  
    ↪values})
```

```
[ ]: X_test_processed
```

```
[ ]: 

|     | raw_data                                          | label |
|-----|---------------------------------------------------|-------|
| 0   | [-8.44707e-05, -8.357633e-05, -9.395467e-05, -... | 1     |
| 1   | [-0.012328528, -0.021660486, -0.024156429, -0...  | 8     |
| 2   | [-0.012697448, -0.017050935, -0.017266486, -0...  | 0     |
| 3   | [0.008832167, 0.0115452595, 0.011243898, 0.011... | 0     |
| 4   | [-0.00033190163, -0.0002754184, -0.00017003811... | 0     |
| ... | ...                                               | ...   |
| 395 | [-0.00018083575, -0.00025713706, -0.0002576629... | 6     |
| 396 | [-8.487413e-05, -0.00027589238, -0.00045009854... | 4     |
| 397 | [-0.008284398, -0.013819212, -0.015761238, -0...  | 8     |
| 398 | [0.009452392, 0.012606097, 0.01237566, 0.01192... | 1     |
| 399 | [-0.00015275163, -0.00018364502, -0.0001810008... | 4     |


```

```
[400 rows x 2 columns]
```

```
[ ]: max_length=17640
```

```
[ ]: def pad_mask (df,max_length):  
    new_padded=[]  
    new_masks = []  
    for i in range (df['raw_data'].shape[0]):  
        length = df['raw_data'][i].shape[0]  
        if length <= max_length:  
            new_add = [0] * (max_length - length)  
            new_one = np.concatenate((df['raw_data'][i],new_add), axis=0)  
            mask = [True]* length + [False]* (max_length - length)  
  
        else:  
            new_one = df['raw_data'][i][: max_length]    ##https://www.geeksforgeeks.  
            ↪org/python-remove-last-k-elements-of-list/#:~:  
            ↪text=%23%20initializing%20list%C2%A0,str(res))  
            mask = [True]*max_length  
            # new_one = np.concatenate((X_train_processed['raw_data'][i],new_add),  
            ↪axis=0)  
            new_padded.append(new_one)  
            new_masks.append(mask)  
    return np.array(new_padded), np.array(new_masks)
```

```
[ ]: X_train_pad_seq, X_train_mask = pad_mask(X_train_processed, max_length)  
X_test_pad_seq, X_test_mask = pad_mask(X_test_processed, max_length)
```

```
[ ]: X_test_pad_seq.shape
```

```
[ ]: (400, 17640)
```

Note - While fitting your model on the augmented data for model 3 you might face Resource exhaust error. One simple hack to avoid that is save the augmented\_train\_data,augment\_y\_train,test\_data and y\_test to Drive or into your local system. Then restart the runtime so that now you can train your model with full RAM capacity. Upload these files again in the new runtime session perform padding and masking and then fit your model.

```
[ ]: padded = Input(shape = (17640,1))  
masked = Input(shape = (17640), dtype = 'bool')  
  
layer = LSTM(50)(padded, mask = masked)  
layer = Dense(128, activation = 'relu', kernel_initializer = 'he_normal')(layer)  
output = Dense(10, activation = 'softmax', kernel_initializer =  
    ↪'glorot_normal')(layer)
```

```
model3 = Model(inputs = [padded,masked], outputs = output)
```

```
[ ]: model3.summary()
```

```
Model: "model_1"
```

```
-----
Layer (type)                 Output Shape          Param #   Connected to
=====
input_2 (InputLayer)         [(None, 17640, 1)]    0         []
input_3 (InputLayer)         [(None, 17640)]       0         []
lstm_1 (LSTM)                 (None, 50)            10400     ['input_2[0][0]',
'input_3[0][0]']
dense_3 (Dense)               (None, 128)           6528      ['lstm_1[0][0]']
dense_4 (Dense)               (None, 10)            1290      ['dense_3[0][0]']
=====
Total params: 18,218
Trainable params: 18,218
Non-trainable params: 0
-----
```

```
[ ]: !rm -rf ./logs/
```

```
[ ]: model3.compile(optimizer = tf.keras.optimizers.Adam(learning_rate = 0.001),
                    loss = 'sparse_categorical_crossentropy')

callback_list = [tf.keras.callbacks.TensorBoard(logdir, histogram_freq=1,
↪write_graph = True),
                 tf.keras.callbacks.EarlyStopping(monitor = 'val_loss', min_delta = 0.05,
↪patience = 5),
                 Metrics([X_test_pad_seq, X_test_mask], y_test)]

model3.fit([X_train_pad_seq, X_train_mask], X_train_processed.label,
↪validation_data = ([X_test_pad_seq, X_test_mask], y_test),
          epochs = 25, callbacks = callback_list, batch_size = 128)
```

```

Epoch 1/25
13/13 [=====] - 3s 145ms/step
val_F1_score: 0.0825
113/113 [=====] - 72s 584ms/step - loss: 2.3030 -
val_loss: 2.3026
Epoch 2/25
13/13 [=====] - 3s 201ms/step
val_F1_score: 0.10749999999999998
113/113 [=====] - 62s 548ms/step - loss: 2.3030 -
val_loss: 2.3027
Epoch 3/25
13/13 [=====] - 2s 144ms/step
val_F1_score: 0.10000000000000002
113/113 [=====] - 59s 523ms/step - loss: 2.3029 -
val_loss: 2.3027
Epoch 4/25
13/13 [=====] - 2s 144ms/step
val_F1_score: 0.10000000000000002
113/113 [=====] - 59s 519ms/step - loss: 2.3028 -
val_loss: 2.3027
Epoch 5/25
13/13 [=====] - 3s 195ms/step
val_F1_score: 0.0975
113/113 [=====] - 60s 529ms/step - loss: 2.3028 -
val_loss: 2.3027
Epoch 6/25
13/13 [=====] - 2s 142ms/step
val_F1_score: 0.10000000000000002
113/113 [=====] - 58s 514ms/step - loss: 2.3028 -
val_loss: 2.3026

```

```
[ ]: <keras.callbacks.History at 0x7a99e00939a0>
```

```
[ ]: %tensorboard --logdir logs/
```

```

Reusing TensorBoard on port 6006 (pid 3892), started 0:13:44 ago. (Use '!kill_
  3892' to kill it.)

```

```
<IPython.core.display.Javascript object>
```

### 1.2.1 4. Data augmentation with spectrogram data

1. use `convert_to_spectrogram` and convert the padded data from train and test data to spectrogram data.
2. The shape of train data will be 14400 x 64 x 35 and shape of test\_data will be 400 x 64 x 35
3. Define the model similar to model 2 and fit the data

```
[ ]: X_train_spectrogram = create_spectrogram(X_train_pad_seq)
```

```
[ ]: X_test_spectrogram = create_spectrogram(X_test_pad_seq)
```

```
[ ]: X_test_spectrogram.shape
```

```
[ ]: (400, 64, 35)
```

```
[ ]: # write the architecture of the model
# print model.summary and make sure that it is following point 2 mentioned above
input = Input(shape = (64,35,))
layer = LSTM(256, return_sequences = True)(input)
layer = tf.math.reduce_mean(layer, axis = -1)
layer = Dense(256, activation = 'relu', kernel_initializer = 'he_normal')(layer)
layer = Dense(128, activation = 'relu', kernel_initializer = 'glorot_normal')(layer)
output = Dense(10, activation = 'softmax', kernel_initializer = 'glorot_normal')(layer)

model4 = Model(inputs = input, outputs = output)
# printing the model summary
model4.summary()
```

Model: "model\_2"

Layer (type)	Output Shape	Param #
input_4 (InputLayer)	[(None, 64, 35)]	0
lstm_2 (LSTM)	(None, 64, 256)	299008
tf.math.reduce_mean_1 (TFOp Lambda)	(None, 64)	0
dense_5 (Dense)	(None, 256)	16640
dense_6 (Dense)	(None, 128)	32896
dense_7 (Dense)	(None, 10)	1290

=====  
Total params: 349,834  
Trainable params: 349,834  
Non-trainable params: 0  
=====

```
[ ]: !rm -rf ./logs/
```

```
[ ]: model4.compile(optimizer = tf.keras.optimizers.Adam(0.001), loss =  
    ↪ 'sparse_categorical_crossentropy')  
callback_list = [tf.keras.callbacks.TensorBoard(logdir, histogram_freq = 1,  
    ↪ write_graph = True),  
    Metrics([X_test_spectrogram, y_test]),  
    tf.keras.callbacks.EarlyStopping(monitor = 'val_loss', min_delta = 0.001,  
    ↪ patience = 3)]  
  
model4.fit(X_train_spectrogram, X_train_processed.label, validation_data =  
    ↪ (X_test_spectrogram, y_test),  
    batch_size = 128, epochs = 50, callbacks = callback_list)
```

Epoch 1/50

4/113 [>...] - ETA: 2s - loss: 0.0984

WARNING:tensorflow:Callback method `on\_train\_batch\_end` is slow compared to the batch time (batch time: 0.0101s vs `on\_train\_batch\_end` time: 0.1024s). Check your callbacks.

13/13 [=====] - 1s 4ms/step

val\_F1\_score: 0.9475

113/113 [=====] - 6s 25ms/step - loss: 0.0695 -

val\_loss: 0.1337

Epoch 2/50

13/13 [=====] - 0s 4ms/step

val\_F1\_score: 0.9675

113/113 [=====] - 2s 15ms/step - loss: 0.0622 -

val\_loss: 0.1041

Epoch 3/50

13/13 [=====] - 0s 3ms/step

val\_F1\_score: 0.965

113/113 [=====] - 1s 13ms/step - loss: 0.0683 -

val\_loss: 0.1081

Epoch 4/50

13/13 [=====] - 0s 3ms/step

val\_F1\_score: 0.965

113/113 [=====] - 1s 12ms/step - loss: 0.0578 -

val\_loss: 0.1204

Epoch 5/50

13/13 [=====] - 0s 3ms/step

val\_F1\_score: 0.9725

113/113 [=====] - 1s 13ms/step - loss: 0.0614 -

val\_loss: 0.0761

Epoch 6/50

13/13 [=====] - 0s 3ms/step

val\_F1\_score: 0.9575

113/113 [=====] - 1s 13ms/step - loss: 0.0619 -

val\_loss: 0.0972

```
Epoch 7/50
13/13 [=====] - 0s 3ms/step
val_F1_score: 0.96
113/113 [=====] - 1s 13ms/step - loss: 0.0671 -
val_loss: 0.0979
Epoch 8/50
13/13 [=====] - 0s 3ms/step
val_F1_score: 0.955
113/113 [=====] - 1s 13ms/step - loss: 0.0642 -
val_loss: 0.1164
```

```
[ ]: <keras.callbacks.History at 0x7a9a74de6140>
```

```
[ ]: %tensorboard --logdir logs/
```

```
Reusing TensorBoard on port 6006 (pid 3892), started 0:19:57 ago. (Use '!kill_
↵3892' to kill it.)
```

```
<IPython.core.display.Javascript object>
```

```
[ ]: from prettytable import PrettyTable
table = PrettyTable()
table.field_names = ['S. No.', 'Model', 'Train Loss', 'Val Loss', 'Val F1_
↵Score']
table.add_row([1, 'Raw Data', 2.3028, 2.3026, 0.1])
table.add_row([2, 'Spectrogram', 0.2740, 0.2633, 0.917])
table.add_row([3, 'Raw Data with Data Augmentation', 2.3028, 2.3026, 0.1])
table.add_row([4, 'Spectrogram with Data Augmentation', 0.0642, 0.1164, 0.955])
print(table)
```

```
+-----+-----+-----+-----+-----+
-----+
| S. No. |           Model           | Train Loss | Val Loss | Val F1
Score |
+-----+-----+-----+-----+-----+
-----+
|  1  |           Raw Data           |  2.3028   |  2.3026   |    0.1
|
|  2  |           Spectrogram         |  0.274    |  0.2633   |    0.917
|
|  3  | Raw Data with Data Augmentation |  2.3028   |  2.3026   |    0.1
|
|  4  | Spectrogram with Data Augmentation |  0.0642   |  0.1164   |    0.955
|
+-----+-----+-----+-----+-----+
-----+
```

**CONCLUSION :** 1. with only raw data model is performing very poorly. Even with the augmented data, val\_F1\_score is not increasing. 2. With spectrogram, val\_f1\_score increased tremen-

dously and with the addition of data augmentation, f1\_score went to 0.955.